



Wissam BENHADDAD

Senior Data Science Consultant | Generative AI Specialist

📍 Paris, France

✉️ wissam.benhaddad.pro@gmail.com

☎️ 06 17 93 45 70

🌐 [wissam-benhaddad-a7a090137](https://www.linkedin.com/in/wissam-benhaddad-a7a090137)

🔗 Wissben

Summary

Forward-thinking **Senior Data Scientist** and **Machine Learning Engineer** adept at using cutting-edge technologies to drive impactful solutions. With **5 years of experience**, I led the development of a **Document Analysis solution** at ByLaw, focusing on harnessing **Large Language Models (LLMs)** to extract valuable insights from complex data. By **fine-tuning** LLM algorithms and integrating them with other technologies, we improved document processing efficiency and accuracy. Beyond my project work, I actively contributed to **R&D** efforts, exploring new applications of **Deep Learning** and making valuable enhancements to internal solutions. Driven by a passion for innovation, I am eager to further leverage my expertise in **GenAI** to drive impactful initiatives and contribute to the evolution of cutting-edge solutions.

Experience

Ekimetrics (Client Assignment- BPCE-Vie), Data Science Consultant

Paris, France

Sept 2025 – present
4 months

- Embedded within the internal Data Science team to shape strategic AI initiatives, translating business challenges into actionable use cases and establishing a comprehensive data strategy roadmap
- Designed and championed a fully AI-driven transformation roadmap for the Financial Department, encompassing data governance enhancements, intelligent chatbot deployment, and automated document analysis pipelines to streamline operational workflows
- Partnered with cross-functional stakeholders to prioritize high-impact AI opportunities, balancing technical feasibility with business value to drive measurable efficiency gains
- **Tech Stack:** Python, Azure, GenAI, Document Processing, Conversational AI

Ekimetrics (Client Assignment), KPI Extraction from Unstructured Documents (Confidential Client)

Paris, France

Mar 2025 – Sept 2025
7 months

- Delivered a scoped, production-ready approach for automated KPI extraction, aligning stakeholders on definitions, acceptance criteria, and measurable quality targets
- Designed a robust document understanding pipeline (RAG, semantic chunking, disambiguation) focused on traceability and consistent outputs
- Built stakeholder trust by translating model behavior into clear decision-ready reporting (failure modes, trade-offs, rollout recommendations)
- **Tech Stack:** Azure OpenAI, Databricks, Document AI, Python

Ekimetrics (Client Assignment - CNP Assurances), GenAI Adoption & Governance — Microsoft Copilot

Paris, France

Nov 2024 – Mar 2025
5 months

- Led discovery workshops to translate business pain points into a ranked GenAI backlog (impact vs. feasibility), enabling fast, outcome-driven delivery
- Delivered and documented 40+ high-value Copilot use cases (summarization, KPI extraction, dashboarding, meeting recaps) with clear guardrails and adoption guidance
- Created a repeatable delivery framework (prompt patterns, validation checklist, governance recommendations) used as reference by both client and internal teams
- Drove enablement by training stakeholders, capturing lessons learned, and improving adoption by making outputs reliable and easy to operationalize
- **Tech Stack:** MS365 Copilot, Copilot Studio, Copilot Agents

Ekimetrics (Innovation Team), Senior Data Scientist — GenAI / Applied Research (LLMs) - Owned the transition from R&D to delivery for LLM fine-tuning, turning ambiguous use-cases into a prioritized roadmap with clear success metrics and shippable milestones - Built an end-to-end fine-tuning playbook (data curation, training recipes, evaluation harness, deployment guidelines) to make results reproducible and production-oriented - Led technical choices (SFT/PEFT, architecture trade-offs, evaluation strategy) and drove reviews with senior stakeholders to de-risk scope, cost, and timeline - Became the go-to advisor for LLM fine-tuning across teams, unblocking projects by pairing deep technical guidance with pragmatic delivery plans Tech Stack: Python, SQL, Databricks, PyTorch, Hugging Face, Docker, Git	Paris, France July 2024 – Dec 2025 1 year 6 months
Ekimetrics (Client Assignment - Groupe Crystal), Customer Segmentation & Scoring - Shipped segmentation and scoring recommendations grounded in rigorous EDA and data quality validation to ensure trustworthy conclusions - Owned delivery from roadmap to execution (planning, sprint rituals, stakeholder reviews), keeping scope tightly aligned to business decisions - Communicated actionable outcomes (what to do next, expected trade-offs, risks) to both technical and non-technical stakeholders Tech Stack: Python, SQL, Databricks, Scikit-learn, Azure, Jira	Paris, France Aug 2024 – Jan 2025 6 months
ByLaw, Senior Data Scientist - Built and deployed a production document extraction solution using multi-modal transformers to automate payslip understanding for compensation workflows - Productionized LLM-based extraction across heterogeneous documents (contracts, legal and financial reports) with focus on reliability and measurable quality - Strengthened the MLOps pipeline through code reviews, bug fixes, and deployment hardening to ensure smooth releases and stable inference - Deployed and operated models on AWS SageMaker for real-time inference, improving end-to-end processing efficiency and maintainability - Contributed to team scaling by interviewing and hiring data scientists, and mentoring through reviews and technical guidance Tech Stack: Python, SQL, AWS SageMaker, PyTorch, Hugging Face, Docker, Git	Paris, France June 2023 – July 2024 1 year 2 months
Qwarry, Lead Data Scientist - Explored retail-media-based advertising offers, leading to innovative solutions to enhance company products - Developed advanced ML models and integrated them into company products and services - Mentored junior data scientists, providing technical guidance and fostering their professional growth - Presented ML research findings to stakeholders, driving informed business decisions Tech Stack: AWS SageMaker, Tableau, Python, Pandas, Jira, Trello, Asana	Paris, France Jan 2023 – June 2023 6 months

Qwarry, Data Scientist

- Implemented and enhanced text classification models within AWS SageMaker environment, improving data processing efficiency
- Established an ML CI/CD pipeline using AWS SageMaker pipelines and GitHub actions, streamlining the deployment process
- Designed and implemented socio-demographic targeting solutions using Machine Learning models, optimizing audience segmentation strategies
- Configured and deployed textual data labeling solutions using LabelStudio, improving data annotation efficiency
- Conducted large-scale web content analysis using PySpark in a cloud environment, extracting actionable insights
- Generated informative dashboards and reports on socio-demographic data, enabling data-driven strategies
- **Tech Stack:** AWS SageMaker, GCP Big Query, PySpark, Python, Pandas, Git, Docker, Jira, Trello

Paris, France
Sept 2020 – Dec 2022
2 years 4 months

Qwarry, Intern Data Scientist

- Developed a machine learning toolbox in Python, streamlining model development processes
- Curated and validated datasets for text classification tasks, ensuring data quality and relevance
- Built and deployed ML models for text classification behind a Django API, enhancing text analysis capabilities
- **Tech Stack:** Keras, TensorFlow, Python, Pandas, Django, SQL

Paris, France
Mar 2020 – Aug 2020
6 months

Education

MS **University Paris Descartes**, Machine Learning for Data Science

Paris, France
Sept 2019 – June 2020

MS **University of Science and Technology**, Intelligent Computer Systems

Algiers, Algeria
Sept 2017 – June 2019

BS **University of Science and Technology**, Computer Science

Algiers, Algeria
Sept 2014 – June 2017

Projects

Speact

Jan 2019 – Dec 2019

Speech recognition project for developing a computer manipulation assistant using intent classification and dialog management

- Developed backend and frontend components for voice-controlled computer assistant
- Implemented NLU pipeline with intent classification and dialog management
- Applied reinforcement learning for smart personal assistant capabilities
- **Tech Stack:** Python, Keras, Pandas, NLTK, Reinforcement Learning

El-Muaiyen

Jan 2018 – Dec 2018

Web application for managing a historical dictionary of the Arabic language

- Built full-stack web application with modern frontend and RESTful backend
- Implemented search functionality using Elasticsearch for efficient text retrieval
- Deployed on Google Cloud Platform using Docker and Kubernetes
- **Tech Stack:** Vue.js, Vuetify, Django, Django Rest Framework, Docker, Kubernetes, GCP, NLTK, Elasticsearch

Skills

Programming Languages: Python, SQL, R, C/C++, Java, JavaScript

GenAI & LLMs: LLM Fine-tuning (PEFT, LoRA, QLoRA), RAG Pipelines, Prompt Engineering, RLHF, LangChain, LlamaIndex, Embedding Models, Model Evaluation

ML/DL Frameworks: PyTorch, TensorFlow, Hugging Face Transformers, Scikit-learn, Keras, NumPy, Pandas, NLTK, spaCy

MLOps & Production: Docker, Kubernetes, MLflow, CI/CD Pipelines, Model Monitoring, GitHub Actions, Prefect, Git

Cloud & Infrastructure: AWS (SageMaker, S3, EC2, Lambda), Azure (OpenAI, ML), GCP (Big Query), Databricks, Terraform

Vector Databases & Search: Pinecone, ChromaDB, FAISS, Elasticsearch, Semantic Search

Data Engineering: PySpark, Big Query, Databricks, SQL, Data Pipelines, ETL

AI Deployment & APIs: FastAPI, Flask, Django, REST APIs, Model Serving, Real-time Inference

Specialized AI: Document AI, Multimodal Transformers, NLU, Text Classification, Computer Vision, Agentic AI

Languages

French: Native

English: Professional working proficiency

Arabic: Native

Interests

- Hiking and outdoor activities
- Football
- Reading and Manga
- Music